

STEPS 2.5 TO 6

PRODUCT IDENTITY THROUGH LAUNCH

Four new prompts slot into the pipeline. They exist so the person running a product through it does not have to make marketing decisions. **The prompts decide. The operator selects.**

UPDATED 15 Aug 2026 NEW STEPS 2.5 · 2.6 · 3L · 4a SOURCE prompts/

THE WHOLE FLOW, ONE LINE

1	2	2.5	2.6	3	3L	3.5	4a	4b	5	6
RESEARCH	PRODUCT BRAIN	IDENTITY	CLAIMS	PRODUCT, IMAGES	LABEL	STORE	ANGLES	PAGE	ADS	LAUNCH

☒ NEW PROMPT ☐ EXISTING STEP, UNCHANGED DECIMAL NUMBERING KEEPS THE OLD STEP PDFS VALID

2.5

NEW

PRODUCT IDENTITY

prompts/step-2.5-product-identity.md

Gives the positioning a name, a page title and a label. Runs after the Product Brain and before images. **Nothing downstream can start without it.** You cannot photograph a label that does not exist, and you cannot write a product page without knowing what the product is called.

PASTE IN, IN THIS ORDER

- 01-avatars.md
- 03-solution-landscape.md
- 05-positioning.md
- The reference listing URL and the supplier listing URL
- Any product images you already have, including the competitor's
- The prompt, last

- **1. Branding gate.** Does this product need a label at all. NO for a finished product with no surface, PARTIAL when someone else's mark has to come off, YES for white label. Everything after depends on the answer.
- **2. Regulated label elements.** Every legally required element, or a plain statement that it must be confirmed with the supplier. Never a guess.
- **3. Product name shortlist.** 20 names cut to 5, sitting on our angle rather than the incumbent's, with the 15 cuts shown and their reasons.
- **4. Product page titles.** 8 USP led titles, outcome first, under 70 characters where possible, never past 100.
- **5. Label copy block.** The exact front and back text, not a description of it.
- **6. Label design brief.** Dieline, hierarchy, palette, what must survive a thumbnail, where [BRAND] sits.
- **7. What is still blocking.** Every [SPEC NEEDED], [CONFIRM WITH SUPPLIER] and [CHECK] gathered into one list.

OUT products/<slug>/brain/06-identity.md

THE NAME GOES ON THE BOTTLE. THE TITLE GOES ON THE STORE.

Two different jobs that keep getting confused. The name has to be short enough to print. The page title is USP led, carries the end state, and can be long. A twelve word USP does not fit on a label.

STOP AND ASK

- The supplier has no label template or dieline.
- The product is in a regulated category and the supplier will not confirm requirements in writing.
- Every good name is taken.
- 05-positioning.md contains a blocking verification. Do not name a product around an unconfirmed USP.

2.6

NEW

CLAIMS LEDGER

prompts/step-2.6-claims-ledger.md

The one document a copywriter needs in front of them. Built so someone handling twelve products at once can write a page from this page alone, without opening a single brain document. Doctrine: **be aggressive, do not invent.**

CATEGORY PURPOSE IS AUTOMATIC.

Knee pads stop knee pain on the job. Gnat spray gets rid of gnats and spider mites. A dust mask keeps dust out of your lungs. That is what the product is, not a hypothesis. Never withhold it, never soften it, never list it as needing evidence. We source products that do the job.

PASTE IN

- 1 All five brain docs, 01 to 05
- 2 06-identity.md
- 3 Product images and the supplier listing
- 4 The prompt, last

THE ONLY REAL CONSTRAINT

- **Fabricated evidence.** Studies, percentages, certifications, test results, award claims, customer counts that do not exist.
- It is about invented proof, never about strength of language.
- The review corpus is a **language bank, not a ceiling.** A three star review is one person's experience, not a limit.

WHAT IT PRODUCES, FIVE SECTIONS

- **1. Cleared to claim.** The big one, in the copywriter's voice, ready to lift. What the product is for, everything it works on, the outcomes, against the alternatives, the mechanism.
- **2. Confirm with sourcing.** A parallel supplier to-do, one line each, ready to paste into an email. Often empty. **Never a gate on building.**
- **3. Off limits.** Fabricated evidence, plus genuinely regulated claim types. Nothing else. A category purpose claim never goes here.
- **4. Mandatory disclosures.** Pulled from 06-identity.md section 2. Usually none.
- **5. Verbatim bank.** The 10 to 15 strongest customer quotes, typos intact, sorted by hooks, pain, transformation, objection handling.

BANNED VOCABULARY, ANYWHERE IN THE OUTPUT

~~may help~~

~~can help~~

~~designed to~~

~~formulated to~~

~~intended to~~

~~aims to~~

~~works to~~

~~can assist~~

~~supports~~

~~promotes~~

~~for many users~~

~~results may vary~~

~~in some cases~~

~~up to~~

OUT products/<slug>/brain/07-claims.md ONE PAGE, TWO AT MOST

THE PRE-SHIP CHECK, TWO QUESTIONS

- 1 Does anything cite a statistic, study, certification or test result that does not exist?
- 2 Is every legally required disclosure present?

Clean on both, ship it. Strength of language is a creative decision, not a compliance one.

Source the product and gather the imagery. One rule bites here every time: **open the source images before using them**. Six of the first seven products arrived carrying somebody else's brand.

OUT `products/<slug>/images/` `products/<slug>/label/`

3.L

NEW SUB-STEP

LABEL DESIGN

`prompts/step-3-label-design.md`

Puts our label on the competitor's product shape. Runs only when `06-identity.md` section 1 answered YES or PARTIAL. Products with no label surface skip it entirely.

THE REFERENCE IMAGE IS GEOMETRY ONLY.

The bottle shape, the cap, the proportions, the lighting. The label area gets replaced, and every trace of the original brand comes off, including the cap, the shoulder, the barcode and any embossing.

PASTE IN

- 1 The competitor or supplier reference image
- 2 `06-identity.md` sections 2, 5 and 6
- 3 The prompt

GUARDRAILS IN EVERY GENERATION PROMPT

- Name the real anatomy of the container. Generation models mangle physical products by default.
- Rule out props from the reference. Fruit and cloths imply claims we did not make.
- Never ask for small legal text as readable type. Ask for a plain block, typeset it properly later.

ACCEPTANCE CHECKLIST, RUN AGAINST THE RESULT

- 1 Is every trace of the original brand gone, including cap, shoulder and barcode?
- 2 Is the container shape still the real product's shape?
- 3 Are `[BRAND]` and `[PRODUCT]` visible as tokens, in the right places, with room for a longer word?
- 4 Is the front readable shrunk to 200 pixels?
- 5 Is anything on it a claim we are not entitled to make, checked against `07-claims.md` ?

OUT `label/label-on-product.jpg` `label/label-flat.jpg`
`label/supplier-link.txt`

3.5

SKILL

Generate at least four variations before choosing. `[BRAND]` stays a visible token until the store is assigned at 3.5, then the same layout is regenerated with the real name. The design work does not repeat. That is the entire reason the token exists.

STORE ASSIGNMENT

skills/product-intake

Files a batch of loose product folders into the repo and assigns each one to an umbrella store, then runs the readiness gate. **Do not file products by hand.** This is where `[BRAND]` resolves.

THE RULES IT ENFORCES

- Stores group by **buyer, not category**. One brand holds two products only if one person buys both.
- A new store needs both triggers: 3 or more products that do not fit, **and** one store already profitable.

OUT `products/<slug>/brief.md` PLUS THE REGISTRY ROW

4a

NEW

PAGE ANGLES

prompts/step-4a-page-angles.md

Runs after the store is assigned and before anyone opens a template. The first two pages built off this system were not usable out of the box. The copy was not wrong. Roughly fifteen judgement calls per page were made in someone's head and recorded nowhere. **4a moves those calls into a document and hands the operator a menu.**

BREADTH IS USUALLY THE PITCH.

The first Electric Scissors page spent all ten reasons on one use case, the recycling bin, for a product with five distinct buyers and six real jobs. Narrow a page on purpose, not by accident.

PASTE IN

- 1 All brain docs, 01 to 05
- 2 07-claims.md
- 3 brief.md, which names the store and the offer
- 4 The prompt, last

STRUCTURE OF THE SHEET

- 1. **Awareness read** and the comparison test
- 2. **Product-aware angles.** 6 to 8 concepts for listicle-1
- 3. **Solution-aware angles.** 2 to 4 for listicle-comparison-above-fold, only if substantiated
- 4. **The default**, marked, for an operator who does not want to choose
- 5. **What to run second**
- **Do not build.** Obvious angles the research rules out

A COMPARISON PAGE IS SUBSTANTIATED ONLY WHEN ALL THREE ARE TRUE

- 1 Three or more avatars sit at Solution Aware or beyond.
- 2 The research holds **real failure evidence for both competitor columns**, not just their marketing claims. A table with an empty failure column is three columns of our own opinion.
- 3 Choosing between alternatives is the buyer's **primary objection**. Buyers defecting from a manual method are a defection story, and that is a product-aware listicle.

WHAT EVERY CONCEPT CARRIES, BECAUSE THIS IS WHAT LISTICLE-BUILDER GATHERS

- **H1**, publishable as written, leading with a real number that matches the item count
- **The angle**, one sentence, the single idea the page argues from headline to close
- **The avatar**, named, with the line saying what they are actually chasing
- **The offer**, from brief.md, or [VERIFY PRICE LIVE] if the price is not set
- **Proof available.** The specific voice bank quotes to spend, or NONE and the page gets no proof furniture
- **Customer phrases**, verbatim, three to six
- **The item list**, 5 to 7, and **images needed**, flagging any that do not exist

THE ITEM LIST, THE PART THAT GETS DONE BADLY

- **5 to 7 items, default 6.** If only five earn their place, write five. A filler item is visible to the reader.
- **The last item is always the guarantee or warranty.** Every page, no exceptions. It is the trust element, it sits after the turn, and it counts in the H1's number. A 7 item page is 6 arguments plus the guarantee.
- **The headline is a contract, so pay it off explicitly.** If the H1 promises "the one they keep it for", one item has to visibly be that one, named as such in its heading.
- **Every headline is help-led.** It states a job the product does for the buyer. The mechanism is demoted to a kicker underneath.
- **Sequence them as an argument.** If the order could be shuffled without loss, it is a grab bag.
- **The product enters late**, straight after the list, as the conclusion the items earned. It is the turn, not an item. The guarantee is the final item.
- One idea per item. On a broad page every avatar appears in at least one item. Bodies cap at roughly 60 words.

It gets the dead undercoat out in the bath instead of on your floor.
Soap free concentrate, releases what the brush cannot reach.

Not this: Contains a soap free deshedding concentrate. Spec-led, the brand admiring itself.

Not this either: Most brushes only reach the surface hair. Observation-led, the reader does the work.

OUT `products/<slug>/brain/08-angles.md`

STOP AND ASK

- Fewer than three angles survive. Given how wide section 1 of the ledger should be, the product is not ready.
- The comparison page is unsubstantiated and the operator specifically wants one.
- Every angle depends on something in OFF LIMITS.

4b

EXISTS

BUILD THE PAGE

templates/landing-pages/ · skills/listicle-builder

The operator makes a selection, not a marketing decision.

08-angles.md Operator picks a row	listicle-builder Hand it that concept block plus 07-claims.md	The copy Headline, intro, 5 to 7 argued items, the product turn, the close
The template Named on the row. Drop it in, deploy, report the full URL		

If the builder has to go back into `05-positioning.md` to make a decision, 4a was not specific enough. That is the test of whether it worked.

HOUSE RULES THE COPY CARRIES

- **Intro length follows awareness.** Product aware or solution aware gets 40 to 70 words, set the frame and hand over to the list. Problem aware or unaware gets 80 to 150, open inside their experience of the problem. Either way, do not name the product in the intro.
- **The final item is the guarantee,** written help-led with the terms as the kicker, and a body that names the real test the buyer should run rather than restating the terms.
- **The close does not repeat the risk reversal,** because the guarantee now has its own item. 45 to 120 words, the offer and one call to action.
- Item bodies stay under 60 words. Real photos at 3:2, four gallery thumbnails, the first screen belongs to the hook.

THE THREE TEMPLATES

- **listicle-1**. The workhorse advertorial. Route per lander in one Next.js app, proven at roughly 40 live routes. Reason bodies under 60 words, real photos at 3:2, four gallery thumbnails, the first screen belongs to the hook.
- **listicle-comparison-above-fold**. Same DNA with a 3-way comparison table directly under the article header. For categories where the alternative is the real objection.
- **ryze-hero-1**. A DTC hero and homepage shell. Not an advertorial, and no live deployment yet.

OUT

products/<slug>/lander/

DEPLOYED, FULL URL REPORTED

MAKE THE ADS

skills/ · brain/creative-system.md

Read `brain/creative-system.md` before any creative work. Do not re-derive the vocabulary. Read that product's `07-claims.md` before writing a single line of ad copy.

static-to-video-multiplier

One winning static becomes 7 to 8 shootable scripts, each with its own hook

video-row-builder

Each script becomes a Notion row an editor can execute without asking a question

higgsfield-generate

The assets themselves, image, video, audio

rapid-ads-workflow

Splits the batch per Facebook Page and per CTA, writes the bulk-copy CSV

In rapid-ads-workflow, "the page" means the Facebook Page identity, not the landing page. Per-launch work is writing the JSON config. The script does not change.

LAUNCH AND TEST

brain/creative-testing-methodology.md

The canonical methodology governs from here: the creative hierarchy, the four modes of testing, single-variable discipline, hypothesis-first design and the pre-launch checklist.

- **A/B testing on landers is Intelligems**, which needs `igId` appended. Shoplift is Shopify-side only.
- Always send the full URL, never a bare slug.
- Do not strip vendor scripts. Pixel and tracking snippets are atomic.
- What is learned goes to `lessons/` in the same session, one fact per file, plus its line in `lessons/INDEX.md`. A lesson that lives only in local memory never reaches the other machine.

WHAT A FINISHED PRODUCT FOLDER HOLDS

Blue is what the new prompts add.

```
products/<slug>/
  brief.md                store assignment, blockers, next action
  brain/
    01-avatars.md
    02-voice-bank.md
    03-solution-landscape.md
    04-gaps.md
    05-positioning.md
    06-identity.md        name, product titles, label copy and brief
    07-claims.md          how far we can go, and the two places we cannot
    08-angles.md          the page menu the operator picks from
  research/
  images/
  label/                  white label only
  lander/                 once a page is built
```

RULES EVERY PROMPT CARRIES

Repeated in each file because these are the ones that get dropped.

- No em-dashes, anywhere.
- Never invent a spec, a statistic, a study, a certification or a volume number. Write a [BRACKETED PLACEHOLDER] instead. **An honest hole beats a filled guess.**
- [BRAND] and [PRODUCT] are never filled in by the model.
- Verbatim customer quotes keep their original spelling and typos.
- The operator picks names and angles. The model shortlists.
- Anything unresolved goes in a single "what is still blocking" list, so there is one place to look.
- Test a prompt on a product where the answer is already known. electric-scissors and knee-crawler both have finished pages built by hand.

Tism Comm Project Steps 1, 2, 3, 4b, 5 and 6 are unchanged

Full detail lives in prompts/ in the repo 15 August 2026